

COM SCI 111 (Operating System Principles)

October 13, 2022

1 9.22 0th

```
$ ls -l big
-rw-r--r-- 1 eggert faculty 9223382036854775000 Sep 22 11:31 big
$ grep x big
$ time grep x big
real 0m0.009s
```

- grep scans at 10^{21} bytes/s (8 Zb/s) searching for an existent line in the file
- so it has to be doing smth other than sequentially searching thru the file
- “grep cheats”
 - “i know it cheats cuz i put the code in it” - eggert 2022
- “i hope you gain the intuition to know when is it ok to cheat and when is it not”
- paul eggert
- <https://web.cs.ucla.edu/classes/fall22/cs111>
- prereqs: cs32 <- c++, algs, data structures / cs118 networking 33 <- computer org, machine code, etc.
|-> cs111 - ... 35l <- shell, python, scripting, ... / ...
- cs131 programming languages
- cs151b architecture
- textbooks
 - ad os 3ep (2018)
 - sk systems (2009)

1.1 whats a system

- *oxford english dictionary* (1928)
 - i. an organized or connected group of objects
 - ii. a set of principles, etc.; a scheme, method
 - from ancient greek $\sigma\iota\sigma\tau\eta\mu\alpha$ (roots, “set up w”)
- *principle of computer science design: an introduction* (2009)
 - smth that ops in an env and the boundary betw the 2 is called the interface
 - interface v important
 - system built from a lot of subsystems
 - standard design: decompose big system
 - often time its useful to have multiple views of the same system
 - sometimes we can look at the system from a diff viewpoint and come up w a completely different set of subsystems

1.2 operating system

- *american heritage dictionary* (2000)
 - software designed to control the hardware of a specific data processing system in order to allow users and application programs to make use of it
 - claims os is very system dependent, not true
- *encarta* (2007)
 - master control program in a computer
- *wikipedia* (2016/8)
 - system software that manages computer, hardware, software resources, and provides common services for computer programs

1.3 goals of an operating system

- protection (of apps, data, . . .)
- performance (from users view)
- utilization (from check-writers view)
- robustness
 - does this operating system do well when given problems out of the ordinary
- flexibility
- simplicity / ease of use
- portability w diff hardware
- scalability
- safety

1.4 what are our main tools?

- abstraction + modularity
 - abstraction: look at the big pic of the system from a particular viewpoint and discard details to understand everything abt that aspect of the system
 - modularity: splitting up a big problem into little problems
 - since the diff in writing a program scales worse than linearly,
 - interface vs implementation
 - mechanisms vs policy
 - policy: high level concept in which u say what u want
 - “i want my interactive processes to have higher priority than background batches”
 - mechanism: how u actually get that stuff to work
- measurements + monitoring
 - measurements: measure how well ur system is working
 - performance + correctness

- monitoring: monitor measurements, do smth w them
- *operating systems: three easy pieces*: main problems in os
 - virtualization: how to build efficient + effective virtual systems
 - concurrency: interacting, simultaneous tasks
 - persistence: data survive failures in hardware, software
 - more
 - security

1.5 a bad os interface

```
char* readline(int fd);
```

- assumes infinite resources

2 9.23 0f

- kernel: lowest level of the os
 - decides what resources are available to apps (thus protecting the system)
 - provides layer of abstraction so that apps dont have to deal w hardware
 - todo: see notes.md.old

2.1 kernel modules

- are pieces of code that can be loaded and unloaded into the kernel upon demand
- they extend the functionality of the kernel without the need to reboot the system
- why not just add all the new functionalities into the kernel image
 - would be bloated
 - security implications
- modules are stored in `/usr/lib/modules/kernel_release`
- to see what kernel modules are currently loaded use

```
lsmod
```

```
cat /proc/modules
```

- example

```
#include <linux/module.h> /* needed by all modules */
```

```
#include <linux/kernel.h> /* needed for KERN_INFO */
```

```
int init_module(void) {
    printk(KERN_INFO "hello world 1.\n");
    return 0;
}

void cleanup_module(void) {
    printk(KERN_INFO "goodbye world 1.\n");
}
```

– printk

– was not meant to communicate info to user

– to load a module: `sudo insmod <module_name> [args]`

```
sudo insmod proc_count.ko
```

– to unload a module: `sudo rmmod <module_name>`

```
sudo rmmod proc_count
```

3 9.27 1t

3.1 simple os architecture

– todo: insert pic

– user mode code: apps, libraries / c-lib

– kernel mode: you can execute instructions here

– e.g. `inb`

– kernel-app interface: “imaginary instructions” that are not real

– `x84-64` instructions `ret`: “`rip = *(--rsp)`” in user mode

– `syscall 35`: “`os[35]()`” in kernel mode

– `syscalls` cannot be called in user mode, kernel will figure it out

```
char *readline(int fd)
```

– `fd`

– `0` - `stdin`

– `1` - `stdout`

– `2` - `stderr`

– `>2` - other files

– points to a newly allocated buffer, e.g. “`ab\0xyz\n`”

– problems

– unbounded cost

– allocates memory for the application

– large potential for memory leaks & bad pointers

- problems with giving the job to kernel
 - force same memory mgmt on all apps
 - syscall overhead

```
ssize_t read(int fd, void *buf, size_t bufsize);
```

- it is now the applications job to figure allocation out, kernel doesn't care
- there is now a limit on read size
- comes at a cost: a program that counts the number of lines in a file becomes more complicated although it has nice properties for the kernel

3.2 problems with designing operating systems

- waterbed effect (tradeoffs): general problem in systems
- incommensurate scaling
 - economies of scale (pin factory)
 - if u just want a few pins then just go to local blacksmith
 - if u want 10mil pins u should create a pin factory
 - diseconomies of scale (star network)
- emergent properties (as u scale)
 - qualitative instead of quantitative changes
 - ucla in school network used for music pirating
- propagation of effects
 - 2 features that independently work may not work when combined
 - e.g. msft invented shift-jis to encode japanese characters
 - 2-byte encoding w top bit on = 2^{15} characters
 - other feature: file names `C:\abc\def\ghi.txt`
 - combination doesn't work because 2nd byte of shift-jis character can be anything (could just happen to be '\')
 - fixed by moving into kernel (complicating the os)
- complexity
 - moore's law

3.3 app: count words in a file

- todo: add pic
- power button = count number of words and put it on screen
- historically called a standalone program
 - operates without benefits of an os
- modern desktop
 - todo: add pic

3.4 uefi: unified extensible firmware interface

- os-independent way to boot
- efi format for **bootloaders**
 - bootloader: find a program in 2ndary storage, copy it into ram and `jmp` to it
 - can be chained, not uncommon to see 3 or 4 chained bootloaders
 - for portability and stuff
 - each bootloader can be more complicated than the next one until one that knows how to boot linux
- guid: globally unique identifier
 - for partitions
 - 128-bit integer
- guid partion table
- uefi boot manager (in firmware)
 - read only but configurable via parameters in some sort of nvram (nonvolatile ram)
- can read gpt tables
- can access files in vfat format etc
- can run code in efi format
- 6 phases

3.5 coreboot

- more hardware-specific / lower level
- no large boot manager (but a small one)
- 4 phases
 - test rom (find out where it is) + flash / disk
 - test ram, early initialization of chipset
 - ram stage: init cpu, chipset, motherboard, devices, etc
 - load payload

3.6 intel core i3-9100

- supports an older, simpler way of booting
- 6 mib l3 cache, 3.6 ghz + 4 gib ddr3 sdram + 1 tb flash sata + intel uhd graphic
- sata: serial ata
 - 7-conductor connector
 - ata (pata): advance technology attachment (16-bit connector in parallel)
 - sequentization is the bane of parallel
 - before ata = ide: integrated drive electronics (western digital, 1986)
- x86 boot procedure
 - 1 mib of physical ram cpu can access

- cpu starts off in real mode = no virtual memory
- initial program counter ip points to rom @ $0xffff0 = 2^{20} - 16$
- program in rom = bios: basic input output system
 - “horrible misdesign”
 - originally user apps ran in real mode, calling subroutines in bios
 - only protection is rom being read only
 - library + kernel + mini os basically
 - bios tries to do the 4 steps in coreboot
 - run in cache only mode
 - step 4: looks for a device containing a particular bit pattern in its first 512 bytes (sector) (mbr: master boot record, 446 bytes of x86 machine code, 64 bytes of partition table (list of 4 pieces of drive, takes role of gpt), 2 bytes of signature $0x55\ 0xaa$ or $0xaa55$)

4 9.29 1th

4.1 less modular os

- one happy program
- uses function calls
- count lines in a file only via function calls + machine instructions at lowest level
- last time we got 426 of bytes of machine code into $0x7c00$
- file is in flash drive
- bootloader reads word count program, say 20 sectors (of 512 bytes), into ram, say $0x1000$

```
static void read_ide_sector(long secno, char *addr) {
    while ((inb(0x1f7) & 0xc0) != 0x40) continue;
    outb(0x1f2, 1);
    outb(0x1f3, secno);
    outb(0x1f4, secno >> 8);
    outb(0x1f5, secno >> 16);
    outb(0x1f6, secno >> 24);
    outb(0x1f7, 0x20);
    while ((inb(0x1f7) & 0xc0) != 0x40) continue;
    insl(0x1f0, addr, 128);
}
```

- inb reads from io registers
- $0x1f7$: status + control
- first while loop waits for drive to be ready
- insl: “insert long”, read 128 words (512 bytes) from $1f0$ to addr


```
static unsigned char inb(unsigned short port) {
    unsigned char data;
    asm volatile("inb %0, %1" : "=a" (data) : "dN" (port));
    return data;
}
```

4.2 i/o

- programmed i/o (pio): special instructions for io
- memory-mapped i/o: more popular, use ordinary load / store instructions `mov_ _`
- intel: flash drive uses pio, monitor uses memory-mapped

```
void write_integer_to_console(long n) {
    unsigned short *p = (unsigned short*) 0xb8014;
    while (n) {
        *p-- = (7 << 8) | ('0' + n % 10);
        n /= 10;
    }
}
```

- at `0xb8000`, 2 bytes per char, 1st byte being format (7: gray on black), 2nd byte being ascii, 80×25 grid
- code above prints

```
static int isalpha(int x) { return 'a' <= x && x <= 'z' || 'A' <= x && x <= 'Z'; }
```

```
void main(void) {
    long words = 0;
    bool in_word = false;
    long secno = 100'000;
    for (;;) ++secno) {
        char buf[512];
        read_ide_sector(secno, buf);
        for (int j = 0; j < 512; ++j) {
            if (!buf[j]) {
                write_integer_to_console(words + in_word);
                while (true);
            }
            bool this_alpha = isalpha((unsigned char) buf[j]);
            words += in_word & ~this_alpha;
            in_word = this_alpha;
        }
    }
}
```

- performance issues

- read 1 sector at a time: change to 255 sectors
- bus contention (controller \leftrightarrow cpu, cpu \leftrightarrow ram): dma (direct memory access), controller \leftrightarrow ram, cpu sends instruction to controller telling where to copy data to
- cpu and controller doing work disjointly: double buffering—cpu issues command to read sector $n + 1$ then cpu counts words in sector n = overlapping work

5 10.4 2t

5.1 how not to have an os standalone program

- standalone program
 - pro: embedded systems
 - con: double buffering api is more complex
 - con: multitasking

```
read_id_sector(...) {
    while (inb(...) & ...) {
        --> schedule(); <--
    }
}
```
- con: dma (direct memory access)
- want separation of concern, guy writing word count program only focuses on counting words
- need better modularity

5.2 what's wrong with function call modularity for apps

- too much pain to change implementation
- too much pain to reuse parts of os in other apps
- too much pain to run simultaneous apps
- too much pain to recover from faults
- not enough insulation between apps

5.3 advantages of modularity

- assume bugs $\propto (k)\text{loc}$ ((thousand) lines of code)
- assume cost to find & fix a bug $\propto \text{kloc}$
- cost to fix all bugs: $O(\text{kloc}^2)$
 - not accurate: more bugs may appear, bugs appear superlinearly to kloc, bugs get harder to fix, etc
- modular: assume k modules, bugs are isolated $\rightarrow O\left(k \cdot \left(\frac{\text{kloc}}{k}\right)^2\right) = O\left(\frac{\text{kloc}}{k}\right)$

5.4 how can we tell whether our modularization is good or bad?

- performance
 - usually willing to sacrifice, say, 5% to 10%
- robustness: tolerance of faults / failures / errors
 - error: mistake in ur head
 - fault: potential problem in code
 - failure: observable behavior that is wrong
 - e.g. dereferencing a null pointer
- neutrality / flexibility / lack of assumptions
- simplicity (of use / to learn)

5.5 mechanism for modularity

0. no modularity
1. function call modularity (call and return instructions)
 - callee can modify
 - callee can loop forever
 - callee can overflow the stack
 - callee can mess w wrong devices
 - callee can execute invalid instruction
 - soft modularity: if callee and caller are both well-behaved / bug-free
 - we want hard modularity: barrier

5.6 3 fundamental system abstractions

1. memory
 - `write(addr, value)`
 - `value = read(addr)`
 - ram, secondary storage
 - issues: throughput, latency, word size, volatility, coherence with caches
2. interpreter
 - `ip` instruction pointer + `ep` environment pointer + repertoire (instruction set)
 - in x86: `rip` and `rsp`
 - above for normal execution
 - interrupts: when normal execution is disturbed
3. link
 - think of 2 modules as different devices, send signals from one to another thru a link

5.7 2 major ways to get hard modularity

1. client / service

- client and server w a link
- if client wants work done, send signal thru link and wait
- client

```
send(factn, {"!", 5});
receive(factn, response);
if response.code == ok:
    print(response.val)
else:
    print("error", response.err)
```

- service

```
for (;;) {
    receive(factn, req);
    if (req.op.code == "!") {
        for (int i = 1; i <= req.n; ++i) {
            a *= i;
        }
        response = { "ok", a };
    } else {
        response = { "bad", 0 };
    }
    send(factn, response);
}
```

- pro: limits error propagation
 - client / server dying doesnt affect the other
- pro: no shared state
- pro: even if service loops forever, client can still make progress
 - timeout on receive
- con: more setup hassle (configuration overhead)
- con: more resources (in the simplest case, need 2 cpus)
- con: latency / thruput / reliability (thru network)

2. virtualization

- using interpreters
- simplest way: write an x86 emulator

```
int epi, eax;

for (;;) {
    char i = *epi++;
```

```
switch (i) {  
  case ...:  
    add(eax, ...);  
}  
}
```

- “client” code: interpreted
- “service” code: functions called by interpreter
- pro: can put whatever checking code for safety
 - e.g. checking pointers before dereferencing
- con: performance (2x - 10x)
- for better speed
 - virtualizable hardware
 - user-mode instructions: `addq`, `call`, `ret`, `jlt`, ...
 - run at full speed in a virtualized (hw interpreter) program
 - kernel-mode instructions: `inb`, `outb`, `insl`, `retl`, ...
 - need them to be rare
 - we need **protected transfer of control** for hard modularity via virtualization
 - `int, 0x80`: interrupt, traps in user mode
 - call kernel w
 - `eax` = syscall number
 - args 1..n: `ebx`, `ecx`, `edx`, `esi`, `edi`, `ebp`
 - hw pushes onto kernel stack
 - `ss`: stack segment
 - `esp`: (user) stack pointer
 - `eflags`
 - `cs`: code segment
 - `cip` (code) stack pointer
 - error code: type of trap
 - trap table: goes to code in kernel

6 10.6 2th

6.1 orthogonality

- processes, races
- how you handle files, stream / network io, processes, should not interfere with each other
- api / interfaces should be
 - simple
 - complete
 - composable

6.2 last time

- INT 0x80 “meta instruction” → pushes onto kernel stack ip and ep, goes into kernel mode
 - kernel code does the syscall in question
 - reti instruction: inverse, pops ip and ep off the stack into register, gets out of kernel mode
 - reti can go to anywhere or run some other process
- seasmnt: x86-64
 - usermode does not use INT instruction
 - uses SYSCALL: uses INT but the idea is that it's faster
 - INT is slower than a function call
 - pushing 6 words instead of 1
 - have to access ram (kernel stack)
 - make sure cache is right (context switch)
 - attempts to streamline, transfer to kernel mode and let kernel figure out what to use and what to restore
 - rax = syscall #
 - args: rdi, rsi, rdx, r10, r8, r9
 - sets rip, etc. to model specific registers
 - one way to implement “syscalls”
- getpid()
- vdso: virtual dynamically-linked shared object
 - ldd /bin/sh
 - libc.so.6 ⇒ /lib64/libc.so.6 shared object code
 - more complicated things such as open()
 - linux-vdso.1 ⇒ kernel memory, readonly for users
 - getpid() in here

6.3 processes

- built from virtualizable processor + os = program in execution in an isolated domain
 - safety
 - simplicity
- processes need to access
 - registers: give cpu 1 process at a time
- cr0 points at page table, maps virtual to physical addresses
- access registers (fast)
- access primary memory: each process has its own page table
- access io devices (less common): syscall, devices do differ
 - storage: flash, hard drive
 - request / response

- random access
 - finite
- stream: keyboard, network
 - spontaneous data generation
 - infinite
- graphics: high performance
 - kind of a mix
- need a set of syscall primitives
 - `fd = open("/usr/lib/libc.so", O_WRONLY | O_NOFOLLOW | ..., 0644)`
 - puts a pointer (file descriptor = index) to file in file descriptor table in process table
 - “opaque handle”
 - `close(12)`
 - `read(12, buf, 512)`
 - will return -1, `errno = EBADF` bad file descriptor
 - has an implicit offset, needs to be stored somewhere
 - stored in file description pointed to by file descriptor
 - which then points to actual file
 - `(sed 1q; sed 's/a/b/;') < file`
 - second sed should start where first stopped
 - so offset shouldn't live in the sed process itself and hence the file description layer between file descriptor and actual file
 - `write(12, buf, 1024)`
 - `n = dup(12) → 15`
 - same file descriptor
 - by convention, file descriptors 0, 1, 2 are `stdin`, `stdout`, `stderr`
 - `int fd0 = dup(0); close(0); int fd1 = open(...);`
- limitations
 - access is sequential
 - potential fix: have a different flavor of `read` for storage devices taking an extra argument
 - orthogonality → `lseek(12, 192308, 0)` position read / write pointer at 192308 from file start (`lseek(fd, offset, whence)`), whence: 1 = from file start `SEEK_SET`, 2 = from current location `SEEK_CUR`, 3 = from file end `SEEK_END`
 - `lseek` read / `lseek` write hurts performance → `pread` / `pwrite`: positioning read / write

6.4 next lecture

- process primitives
 - `fork()`

7 10.11 3t

7.1 file descriptor trouble

- use a fd thats closed (or never opened)

```
#include <errno>
```

```
write(47, "xy", 2)
```

- returns -1, sets errno == EBADF

- open, but resource not available

```
int f = open(...);  
if (f < 0) { error; return; }  
read(f);  
somefun(f);  
write(1, ...);
```

- io error
- end of file (read returns 0)
- errno: lots of ifs

7.2 process api in posix/linux

- pid_t fork(void);
 - #include <sys/wait.h>: typedef int pid_t
 - returns an integer containing a process id
 - returns -1 on failure
 - otherwise returns new process id
 - positive integer
 - returns 0 if in child process
- int execvp(char const *file, char * const *argv);
- pid_t waitpid(pid_t p, int *status, int options);: wait for child process to die and store exit status into an integer in parent, returns p if successfully p to die, otherwise returns -1 (e.g. wrong p, p already dead)

```
pid_t p = fork();  
switch (p) {  
    case -1: return error();  
    case 0:  
        execvp("/bin/date", (char*[]) {"date", "-u", NULL});  
        return error(); // <-- executed in child process
```



```
    default:
        int status;
        if (waitpid(p, &status, 0) < 0) return error(...);
        if (!WIFEXITED(status) || WEXITSTATUS(status) != 0)
            return error(...);
    }
- date.c:

    int main(int argc, char **argv) {...}
- aside on restrict
    - caller not allowed to pass pointers pointing to the same place (todo)

int posix_spawnvp(
    pid_t *restrict pid,
    char const *restrict file,
    posix_spawn_file_actions_t const *restrict acts,
    posix_spawn_attr_t const *restrict attrp,
    char *const *restrict argv,
    char *const *restrict envp
);
- _Noreturn void _exit(int status);
    - send message to parent process

while (fork()) continue;
```

7.3 todo: ??

- todo: hr 2
- copy entry in process table
 - except rax which stores result of system call
 - pid in rax of parent
 - 0 in rax of child
- fork: child = parent *except* for
 - return value of fork
 - pid
 - ppid (parent pid)
 - accumulated execution times
 - file descriptors (file descriptions are shared)
 - file locks (child does not have the lock)
 - pending signals
- exec is opposite: it destroys / replaces program data (stack, heap), registers, signal handlers reset to default

7.4 processes do need to affect each other

- need controlled isolation
- files: simple, straightforward, slow, space, names, hassle
 - race conditions
 - occurs when behavior depends on timing
 - (cat a & cat b) >outfile: nondeterministic
 - (cat >a & cat >b) <infile
 - sorting a big file and using a temp file to store intermediate results: `int tmpfd = open("/tmp/sorttmp", O_WRONLY | O_CREAT | O_TRUNC, 0600);`
 - someone else might have it
 - wait for other sort instance to finish using it
 - still has space where race condition can happen
 - toctou race
 - `acquire_lock("/tmp");`
 - `// begin critical sector`
 - `if (access(...))`
 - `open(...);`
 - `// end critical sector`
 - `release_lock("/tmp");`
 - `fcntl(fd, F_SETLK / F_SETLKW / F_GETLK / F_UNLOC)` system call
 - *voluntary* locks
 - performance problem:
 - `O_EXCL` exclusive flag: fail if file already exists
 - include pid in filename
 - todo: why bad?
 - use random number + `O_EXCL` flag
 - work but somewhat unsatisfactory
 - `O_TMPFILE` flag: create a file somewhere in directory but don't give it a name
 - `int fd = open("/tmp", O_TMPFILE | O_CREAT | ...);`
 - essentially do open and `unlink` atomically
- message-passing
- shared memory: 2 processes not isolated completely, fastest, most dangerous - race conditions
- exit status, signals kill -HUP 1923
- covert channels: e.g. cpu load

8 10.13 3th

8.1 breaking application down

- find data → read input → find words → count words → output
- whos in charge

- lazy evaluation
 - minimize total amt of work done
 - `count_words` in charge
- eager evaluation
 - `read_input` in charge
- mixed evaluation
 - all 3 run simultaneously
 - allows for more parallelism
- pipes
 - control execution of multiple processes without any of the processes being in charge

8.2 what a pipe looks like inside the kernel

- buffer of certain size, depending on the kernel, say 4 KiB
- w/in the buffer there are values `r` / `w`: read / write pointer
- problems with pipes
 - write to a full pipe
 - consider as an error, return `-1`, set `errno = EAGAIN`
 - process blocks (default)
 - process gets `SIGPIPE` signal
 - if no reader
 - read from an empty pipe
 - rare (todo:??)
 - normal
 - if no writers: read returns `0 == EOF`
- pipes are nameless
 - unless named pipe

```
$ mkfifo p
$ ls -l p
prw-rw-rw- ... p

int pipe(int ds[2]);
```
 - to use the pipe, call `fork`
 - `$ A | B` want another shell prompt after this
 - need 3 processes
 - `sh` \rightarrow A, B
 - `cat bigfile | less` \rightarrow q
 - less is gone
 - shell only cares about B fishing

- sh → A, B
- sh → B, A
 - conventionally, this process tree is used
 - performance advantage: more parallelism, child process sets up pipeline
- shell.c

```
pid_t b = fork();
if (b < 0) error();
if (b == 0)
    setup_a_b(...);
else if (waitpid(b, &status, 0) < 0)
    error();
return status;

...
```

```
void setup_a_b(void) {
    int fd[2];
    if (pipe(fd) != 0) error();
    pid_t a = fork();
    if (a < 0) error();
    if (a == 0) {
        dup2(fd[1], STDOUT_FILENO);
        close(fd[0]);
        close(fd[1]);
        execlp("A", "A", NULL);
        error();
    }
    dup2(fd[0], STDIN_FILENO);
    close(fd[0]);
    close(fd[1]);
    execlp("B", "B", NULL);
    error();
}
```

- no exit status for A

8.3 losing power

- small backup power (battery)
- copy cache / ram into secondary storage fast
- how to support this
 - simplify applications by saying kernel will handle everything
 - snapshots each process
 - as if nothing happened in the processes' pov

- relatively expensive (efficiency problem)
- assumes you can snapshot everything
 - often not the case when process is talking to the outside world (e.g. cloud)
 - cannot snapshot everything
 - clock will change (time dependent processes)
- /dev/power
 - tells you the power status
 - make processes inspect /dev/power and do appropriate things
 - each program must be modified to read /dev/power (polling)
- blocking read `char c; fd = open("/dev/power", ...); read(fd, &c, 1);`
 - cant do anything while blocked
 - unless assuming multithreading, but need to notify other threads
- signals
 - rare events
 - localize handling
 - break out of an expensive / complicated loop
 - kill a process
 - timeouts

8.4 signals on linux

- `kill -HUP 19362`
 - send hangup signal to process 19362
 - `if (kill(SIGHUP, 19362) < 0)`
- ordinarily user code can only send signals to processes with the same user id
- user
 - C-c: send SIGINT to every process in the foreground
 - SIGKILL: i want you to die and i dont want you to be able to do anything abt it, cannot be ignored
 - `kill -KILL $$`
 - SIGSTOP
 - SIGUSR1...
 - SIGHUP
- internal errors
 - invalid instruction: SIGILL, illegal instruction signal
 - floating point exception: SIGFPE, not sent anymore, only sent when INT_MIN / -1 or 0 / 0
 - invalid address (bad ptr etc): SIGSEGV, SIGBUS
- io errors
 - SIGIO
 - SIGPIPE: writing to pipe w no readers
- death of a child process

- SIGCHLD: rare, one of subsidiaries has died, prob want to clal waitpid
- alarm clock
 - SIGALRM
- SIGXCPU: sent to process when (just abt to) out of cpu time
- SIGXFSZ: file size
- `#include <signal.h>`
 - `typedef void (*sig_handler_t)(int);` function type, takes `int` and returns `void`
 - `sig_handler_t signal(int signo, sig_handler_t handler)` set signal handler and return previous
- ```
void handle_control_C(int sig) {
 write(1, "?", 1);
}
```
- ```
int main(void) {  
    signal(SIGINT, handle_control_C);  
    ...  
}
```
- signal handling dispatch table in process table
- signal handlers can be called (by default) between any pair of machine instructions
 - instructions are atomic
 - syscalls are atomic
 - others are not: library calls
 - have to be prepared for the handler function to be called anywhere within a library call
 - e.g. `malloc(1000)`
 - `malloc` in handler = disaster
 - posix rule: dont do that, never call `malloc` from signal handler
 - e.g. `printf("?")`
 - may mess up internal data structure
 - reentrant functions: functions that you can call in the middle of their call
 - e.g. `sqrt`
 - !e.g. `malloc, printf`