

COM SCI 111 (Operating System Principles)

September 29, 2022

1 9.22 0th

```
$ ls -l big
-rw-r--r-- 1 eggert faculty 9223382036854775000 Sep 22 11:31 big
$ grep x big
$ time grep x big
real 0m0.009s
```

- grep scans at 10^{21} bytes/s (8 Zb/s) searching for an existent line in the file
- so it has to be doing smth other than sequentially searching thru the file
- “grep cheats”
 - “i know it cheats cuz i put the code in it” - eggert 2022
- “i hope you gain the intuition to know when is it ok to cheat and when is it not”
- paul eggert
- <https://web.cs.ucla.edu/classes/fall22/cs111>
- prereqs: cs32 <- c++, algs, data structurues / cs118 networking 33 <- computer org, machine code, etc.
|-> cs111 - ... 35l <- shell, python, scripting, ... / ...
- cs131 programming languages
- cs151b architecture
- textbooks
 - ad os 3ep (2018)
 - sk systems (2009)

1.1 whats a system

- *oxford english dictionary* (1928)
 - i. an organized or connected group of objects
 - ii. a set of principles, etc.; a scheme, method
 - from ancient greek $\sigma\iota\sigma\tau\eta\mu\alpha$ (roots, “set up w”)
- *principle of computer science design: an introduction* (2009)
 - smth that ops in an env and the boundary betw the 2 is called the interface
 - * interface v important
 - * system built from a lot of subsystems
 - * standard design: decompose big system
 - * often time its useful to have multiple views of the same system
 - sometimes we can look at the system from a diff viewpoint and come up w a completely different set of subsystems

1.2 operating system

- *american heritage dictionary* (2000)
 - software designed to control the hardware of a specific data processing system in order to allow users and application programs to make use of it
 - * claims os is very system dependent, not true
- *encarta* (2007)
 - master control program in a computer
- *wikipedia* (2016/8)
 - system software that manages computer, hardware, software resources, and provides common services for computer programs

1.3 goals of an operating system

- protection (of apps, data, ...)
- performance (from users view)
- utilization (from check-writers view)
- robustness
 - does this operating system do well when given problems out of the ordinary
- flexibility
- simplicity / ease of use
- portability w diff hardware
- scalability
- safety

1.4 what are our main tools?

- abstraction + modularity
 - abstraction: look at the big pic of the system from a particular viewpoint and discard details to understand everything abt that aspect of the system
 - modularity: splitting up a big problem into little problems
 - * since the diff in writing a program scales worse than linearly,
 - interface vs implementation
 - mechanisms vs policy
 - * policy: high level concept in which u say what u want
 - “i want my interactive processes to have higher priority than background batches”
 - * mechanism: how u actually get that stuff to work
- measurements + monitoring
 - measurements: measure how well ur system is working
 - * performance + correctness
 - monitoring: monitor measurements, do smth w them
- *operating systems: three easy pieces*: main problems in os
 - virtualization: how to build efficient + effective virtual systems

- concurrency: interacting, simultaneous tasks
- persistence: data survive failures in hardware, software
- more
 - * security

1.5 a bad os interface

```
char* readline(int fd);
```

- assumes infinite resources

2 9.23 of

- kernel: lowest level of the os
 - decides what resources are available to apps (thus protecting the system)
 - provides layer of abstraction so that apps don't have to deal w hardware
 - TODO: see notes.md.old

2.1 kernel modules

- are pieces of code that can be loaded and unloaded into the kernel upon demand
- they extend the functionality of the kernel without the need to reboot the system
- why not just add all the new functionalities into the kernel image
 - would be bloated
 - security implications
- modules are stored in `/usr/lib/modules/kernel_release`
- to see what kernel modules are currently loaded use

```
lsmod
```

```
cat /proc/modules
```

- example

```
#include <linux/module.h> /* needed by all modules */
```

```
#include <linux/kernel.h> /* needed for KERN_INFO */
```

```
int init_module(void) {  
    printk(KERN_INFO "hello world 1.\n");  
    return 0;  
}
```

```
void cleanup_module(void) {  
    printk(KERN_INFO "goodbye world 1.\n");  
}
```

- `printk`
 - was not meant to communicate info to user
- to load a module: `sudo insmod <module_name> [args]`

```
sudo insmod proc_count.ko
```

- to unload a module: `sudo rmmod <module_name>`

```
sudo rmmod proc_count
```

3 9.27 1t

3.1 simple os architecture

- TODO: insert pic
- user mode code: apps, libraries / c-lib
- kernel mode: you can execute instructions here
 - e.g. `inb`
 - kernel-app interface: “imaginary instructions” that are not real
 - * x84-64 instructions `ret: “rip = *(--rsp)”` in user mode
 - * `syscall 35: “os[35] ()”` in kernel mode
 - syscalls cannot be called in user mode, kernel will figure it out

```
char *readline(int fd)
```

- `fd`
 - 0 - `stdin`
 - 1 - `stdout`
 - 2 - `stderr`
 - >2 - other files
- points to a newly allocated buffer, e.g. `"ab\0xyz\n"`
- problems
 - unbounded cost
 - allocates memory for the application
 - large potential for memory leaks & bad pointers
 - * problems with giving the job to kernel
 - force same memory mgmt on all apps
 - syscall overhead

```
ssize_t read(int fd, void *buf, size_t bufsize);
```

- it is now the applications job to figure allocation out, kernel doesn't care
- there is now a limit on read size
- comes at a cost: a program that counts the number of lines in a file becomes more complicated although it has nice properties for the kernel

3.2 problems with designing operating systems

- waterbed effect (tradeoffs): general problem in systems
- incommensurate scaling
 - economies of scale (pin factory)
 - * if u just want a few pins then just go to local blacksmith
 - * if u want 10mil pins u should create a pin factory
 - diseconomies of scale (star network)
- emergent properties (as u scale)
 - qualitative instead of quantitative changes
 - ucla in school network used for music pirating
- propagation of effects
 - 2 features that independently work may not work when combined
 - e.g. msft invented shift-jis to encode japanese characters
 - * 2-byte encoding w top bit on = 2^{15} characters
 - * other feature: file names `C:\abc\def\ghi.txt`
 - * combination doesnt work because 2nd byte of shift-jis character can be anything (could just happen to be '\')
 - * fixed by moving into kernel (complicating the os)
- complexity
 - moores law

3.3 app: count words in a file

- TODO: add pic
- power button = count number of words and put it on screen
- historically called a standalone program
 - operates without benefits of an os
- modern desktop
 - TODO: add pic

3.4 uefi: unified extensible firmware interface

- os-independent way to boot
- efi format for **bootloaders**
 - bootloader: find a program in 2ndary storage, copy it into ram and `jmp` to it
 - can be chained, not uncommon to see 3 or 4 chained bootloaders
 - * for portability and stuff
 - * each bootloader can be more complicated than the next one until one that knows how to boot linux
- guid: globally unique identifier
 - for partitions
 - 128-bit integer

- guid partition table
- uefi boot manager (in firmware)
 - read only but configurable via parameters in some sort of nvram (nonvolatile ram)
- can read gpt tables
- can access files in vfat format etc
- can run code in efi format
- 6 phases

3.5 coreboot

- more hardware-specific / lower level
- no large boot manager (but a small one)
- 4 phases
 - test rom (find out where it is) + flash / disk
 - test ram, early initialization of chipset
 - ram stage: init cpu, chipset, motherboard, devices, etc
 - load payload

3.6 intel core i3-9100

- supports an older, simpler way of booting
- 6 mib l3 cache, 3.6 ghz + 4 gib ddr3 sdram + 1 tb flash sata + intel uhd graphic
- sata: serial ata
 - 7-conductor connector
 - ata (pata): advance technology attachment (16-bit connector in parallel)
 - * sequentization is the bane of parallel
 - before ata = ide: integrated drive electronics (western digital, 1986)
- x86 boot procedure
 - 1 mib of physical ram cpu can access
 - * cpu starts off in real mode = no virtual memory
 - initial program counter ip points to rom @ $0xffff0 = 2^{20} - 16$
 - program in rom = bios: basic input output system
 - * “horrible misdesign”
 - * originally user apps ran in real mode, calling subroutines in bios
 - only protection is rom being read only
 - * library + kernel + mini os basically
 - * bios tries to do the 4 steps in coreboot
 - run in cache only mode
 - step 4: looks for a device containing a particular bit pattern in its first 512 bytes (sector) (mbr: master boot record, 446 bytes of x86 machine code, 64 bytes of partition table (list of 4 pieces of drive, takes role of gpt), 2 bytes of signature $0x55\ 0xaa$ or $0xaa55$)

4 9.29 1th

4.1 less modular os

- one happy program
- uses function calls
- count lines in a file only via function calls + machine instructions at lowest level
- last time we got 426 of bytes of machine code into 0x7c00
- file is in flash drive
- bootloader reads word count program, say 20 sectors (of 512 bytes), into ram, say 0x1000

```
static void read_id_sector(long secno, char *addr) {
    while ((inb(0x1f7) & 0xc0) != 0x40) continue;
    outb(0x1f2, 1);
    outb(0x1f3, secno);
    outb(0x1f4, secno >> 8);
    outb(0x1f5, secno >> 16);
    outb(0x1f6, secno >> 24);
    outb(0x1f7, 0x20);
    while ((inb(0x1f7) & 0xc0) != 0x40) continue;
    insl(0x1f0, addr, 128);
}
```

- inb reads from io registers
- 0x1f7: status + control
- first while loop waits for drive to be ready
- insl: "insert long", read 128 words (512 bytes) from 1f0 to addr

```
static unsigned char inb(unsigned short port) {
    unsigned char data;
    asm volatile("inb %0, %1" : "=a" (data) : "dN" (port));
    return data;
}
```

4.2 i/o

- programmed i/o (pio): special instructions for io
- memory-mapped i/o: more popular, use ordinary load / store instructions mov_ _
- intel: flash drive uses pio, monitor uses memory-mapped

```
void write_integer_to_console(long n) {
    unsigned short *p = (unsigned short*) 0xb8014;
    while (n) {
        *p-- = (7 << 8) | ('0' + n % 10);
        n /= 10;
    }
}
```


- at 0xb8000, 2 bytes per char, 1st byte being format (7: gray on black), 2nd byte being ascii, 80 × 25 grid
- code above prints

```
static int isalpha(int x) { return 'a' <= x && x <= 'z' || 'A' <= x && x <= 'Z'; }
```

```
void main(void) {  
    long words = 0;  
    bool in_word = false;  
    long secno = 100'000;  
    for (;;) ++secno) {  
        char buf[512];  
        read_ide_sector(secno, buf);  
        for (int j = 0; j < 512; ++j) {  
            if (!buf[j]) {  
                write_integer_to_console(words + in_word);  
                while (true);  
            }  
            bool this_alpha = isalpha((unsigned char) buf[j]);  
            words += in_word & ~this_alpha;  
            in_word = this_alpha;  
        }  
    }  
}
```

- performance issues
 - read 1 sector at a time: change to 255 sectors
 - bus contention (controller ↔ cpu, cpu ↔ ram): dma (direct memory access), controller ↔ ram, cpu sends instruction to controller telling where to copy data to
 - cpu and controller doing work disjointly: double buffering—cpu issues command to read sector $n + 1$ then cpu counts words in sector n = overlapping work